




Google Summer of Code 2025

Proposal



Adding clustering *Searcher classes in statistics package

PERSONAL DETAILS:

Name: Swayam Shah (Sonu)

GitHub: [Sonu0305](#)

Discourse: [Swayam](#)

Telephone: [+91 8291602906](#)

Email: swayamshah66@gmail.com

LinkedIn: [swayamshahh](#)

Country: India

Time Zone: IST (GMT+5:30)

Language: English

PROJECT ABSTRACT:

The **GNU Octave** statistics package **lacks** extensible **clustering search classes** and has an **inefficient**, disabled **KDTree** implementation (**GitHub issue #151**). This project will **implement** **KDTreeSearcher**, **ExhaustiveSearcher**, and **hnswSearcher** classes with **knnsearch** and **rangearch** methods, plus a **createns** helper function. Using a **C++-compiled .oct** library, the **KDTree** **will be optimized** for faster construction and queries, **surpassing** MATLAB compatibility. This will **enhance Octave's** clustering capabilities for **statistical computing**.

OVERVIEW:

GNU Octave is a premier open-source environment for **scientific computing**, celebrated for its **numerical prowess** and status as a free alternative to MATLAB. Its **rich package** ecosystem fuels diverse applications, from academic research to real-world problem-solving, with the **statistics package** playing a pivotal role in **statistical analysis**. Yet, deficiencies in its **clustering tools** reveal an **opportunity** to elevate its capabilities and retain users who might otherwise turn to proprietary alternatives.

This project rectifies this by introducing advanced searcher classes: **KDTreeSearcher** for **fast k-nearest neighbor** searches using a spatial tree to quickly locate the closest points, **ExhaustiveSearcher** for **brute-force** searches ensuring precision across all data, and **hnswSearcher** for **approximate nearest neighbor searches** leveraging hierarchical navigable small world graphs to balance speed and accuracy. These will include **knnsearch** to find the **k closest points** and **rangearch** to retrieve all **points within a given radius**, enabling applications like pattern recognition, outlier detection, density estimation, and geospatial analysis. **Optimized** with compiled C++ code, these features will handle large datasets **efficiently**.

As an **active Octave contributor**, I've engaged in **community discussions**, **assisted users**, and **mastered** its **codebase**, preparing me to **meet user needs**.

Through Google Summer of Code 2025, I'll apply my **Octave**, **C++**, and **clustering expertise**, collaborating with mentors to **enhance Octave's utility**, **broaden** its appeal, and deepen my role in its **vibrant community**. I am **motivated** to contribute and join the community. It's always great to see people use and appreciate your code.

DETAILED PROJECT DESCRIPTION:

I would like to focus on few important areas like:

1. **Lack of Modularity:** The current implementation does not use object-oriented programming (OOP) principles. This makes it difficult to extend or modify the functionality.
2. **Inefficient KDTree Implementation:** The KDTree method is slow and poorly implemented.
3. **Missing hnswSearcher:** The current implementation does not include the Hierarchical Navigable Small World (HNSW) algorithm, which is a state-of-the-art method for approximate nearest neighbor searches.
4. **No Helper Function:** There is no **createns** helper function to create a nearest neighbor searcher object.

I would also like to share the **class diagrams** of proposed implementation of the classes:

| KDTreeSearcher | ExhaustiveSearcher | hnsWSearcher |
|---|---|---|
| - data: matrix - tree: KDTree (C++ object) | - data: matrix | - data: matrix - graph: HNSW (C++ object) |
| + KDTreeSearcher(X) + knnsearch(Y, k): idx, D + rangesearch(Y, r): idx, D | + ExhaustiveSearcher(X) + knnsearch(Y, k): idx, D + rangesearch(Y, r): idx, D | + hnsWSearcher(X) + knnsearch(Y, k): idx, D + rangesearch(Y, r): idx, D |

I have also **researched** about how and decided about on how I would like the **method signatures** to be:

1. KDTreeSearcher:

```
classdef KDTreeSearcher
    properties
        data
        tree
    end
    methods
        function obj = KDTreeSearcher(X)
            % Constructor: Build KDTree from data matrix X
            obj.data = X;
            obj.tree = build_kdtree(X); % C++ function for KDTree construction
        end

        function [idx, D] = knnsearch(obj, Y, k)
            % Find k-nearest neighbors in KDTree
            [idx, D] = kdtree_knnsearch(obj.tree, Y, k); % C++ function for
kNN search
        end

        function [idx, D] = rangesearch(obj, Y, r)
            % Find all neighbors within radius r in KDTree
            [idx, D] = kdtree_rangesearch(obj.tree, Y, r); % C++ function for
range search
        end
    end
end
```

a. Using **KDTreeSearcher** (For Example):

```
X = rand(1000, 10); % 1000 points in 10D space
Y = rand(10, 10); % 10 query points

% Create KDTreeSearcher
searcher = KDTreeSearcher(X);

% Find 5 nearest neighbors
[idx, D] = searcher.knnsearch(Y, 5);

% Find neighbors within radius 0.5
[idx, D] = searcher.rangesearch(Y, 0.5);
```

2. **ExhaustiveSearcher**:

```
classdef ExhaustiveSearcher
    properties
        data
    end
    methods
        function obj = ExhaustiveSearcher(X)
            % Constructor: Store data matrix X
            obj.data = X;
        end

        function [idx, D] = knnsearch(obj, Y, k)
            % Brute-force k-nearest neighbors search
            D = pdist2(obj.data, Y, 'euclidean'); % Use existing pdist2
            function
                [D, idx] = sort(D, 2);
                D = D(:, 1:k);
                idx = idx(:, 1:k);
            end

        function [idx, D] = rangesearch(obj, Y, r)
            % Brute-force range search
            D = pdist2(obj.data, Y, 'euclidean');
            idx = cell(size(Y, 1), 1);
            for i = 1:size(Y, 1)
                idx{i} = find(D(:, i) ≤ r);
            end
        end
    end
end
```

3. hnswSearcher:

```
classdef hnswSearcher
    properties
        data
        graph
    end
    methods
        function obj = hnswSearcher(X)
            % Constructor: Build HNSW graph from data matrix X
            obj.data = X;
            obj.graph = build_hnsw(X); % C++ function for HNSW construction
        end

        function [idx, D] = knnsearch(obj, Y, k)
            % Approximate k-nearest neighbors search using HNSW
            [idx, D] = hnsw_knnsearch(obj.graph, Y, k); % C++ function for
HNSW kNN search
        end

        function [idx, D] = rangesearch(obj, Y, r)
            % Approximate range search using HNSW
            [idx, D] = hnsw_rangesearch(obj.graph, Y, r); % C++ function for
HNSW range search
        end
    end
end
```

a. Using **hnswSearcher** (For Example):

```
X = rand(100000, 100); % 100,000 points in 100D space
Y = rand(10, 100);      % 10 query points

% Create hnswSearcher
searcher = hnswSearcher(X);

% Find 10 approximate nearest neighbors
[idx, D] = searcher.knnsearch(Y, 10);
```

BENEFITS TO OCTAVE COMMUNITY:

1. **Closing the Feature Gap:** Adding **KDTreeSearcher**, **ExhaustiveSearcher**, and **hnswSearcher** aligns Octave's statistics package with **MATLAB**, attracting users needing advanced clustering tools.
2. **Enabling Advanced Research:** Enhanced search capabilities empower data scientists and researchers to **efficiently analyze** large datasets for machine learning and spatial tasks.

3. **Strengthening the Ecosystem:** Improved clustering functionality fosters development of higher-level tools, enriching Octave's open-source ecosystem.

DELIVERABLES:

1. **Searcher Classes:** `KDTreeSearcher`, `ExhaustiveSearcher`, `hnswSearcher` with `knnsearch` and `rangesearch`.
2. **Optimized KDTree:** C++-compiled .oct file **fixing** GitHub issue #151.
3. **createns Function:** Helper to instantiate searchers with options.
4. **Documentation:** Guides and examples for new features.
5. **Tests:** Suite ensuring compatibility and performance.

POSSIBLE DIFFICULTIES:

1. **MATLAB Parity:** Matching MATLAB's behavior might hit edge-case snags. I'll align with its docs and test rigorously.
2. **Octave Integration:** Linking .oct files and new classes could face build issues. My codebase knowledge and early tests will ease this.
3. **C++ Performance:** Optimizing KDTree in C++ for speed and memory may be tricky. I'll use deeper research & mentor input to refine it.

TIMELINE:

1. Once **coding begins**, I will work dedicatedly for **at least 4-5 hours on weekdays** (and **6-8 hours on weekends**) until the completion of the project. After the **submission** of my proposal (**on or before April 8**), I will start **researching more** about the classes and start noting down important aspects of it.
2. During the **start of** the community bonding period, I have my **end-semester examinations** which will start from **May 5 until May 17**. However, I **promise** to begin my work as early as possible to be on track. After the examinations, **I shall put in more time** to make up for any lost time or lag.
3. I have a **2-month long vacation after May 17** and **no other commitments** in hand, so I will be able to **devote ample time** to the project.

| | |
|----------------|--|
| Up till May 8 | Proposal accepted or rejected → Deepen familiarity with Octave's statistics package codebase. → Discuss project details with mentor Andreas Bertsatos, refining scope and priorities . |
| May 8 - June 1 | Community Bonding Period → Set up the development environment with the latest Octave version and C++ tools. → Finalize detailed plan: prioritize <code>KDTreeSearcher</code> , identify potential HNSW challenges . |

| | |
|-----------------------------|--|
| Week 1 & 2 | Coding Officially Begins! <ul style="list-style-type: none"> → Design KDTreeSearcher class using classdef. → Start C++ KDTree implementation (tree-building logic), test basic queries. |
| Week 3 & 4 | <ul style="list-style-type: none"> → Complete C++ KDTree as .oct file, integrate with KDTreeSearcher. → Implement knnsearch and rangearch for KDTreeSearcher, test performance. |
| Week 5 | Buffer Period <ul style="list-style-type: none"> → Fix any KDTree issues, optimize based on initial tests. → Review progress with mentor, adjust remaining tasks. |
| Week 6 & 7 | <ul style="list-style-type: none"> → Develop ExhaustiveSearcher with knnsearch and rangearch, test accuracy. → Implement createns function, validate with both KDTree and Exhaustive methods. |
| Week 8 & 9 | <ul style="list-style-type: none"> → Build hnswSearcher using HNSW algorithm, add methods. → Test hnswSearcher on large datasets, tweak parameters for speed/accuracy. |
| Week 10 | Buffer Period <ul style="list-style-type: none"> → Resolve bugs in hnswSearcher or prior classes. → Ensure all classes work seamlessly with createns, confirm deliverables with mentor. |
| Week 11 & 12 | <ul style="list-style-type: none"> → Write comprehensive documentation and usage examples for all features. → Finalize testing suite (performance, compatibility), prepare submission (code, docs, report). |
| Until Pens Down & Post-GSoC | <ul style="list-style-type: none"> → Address final mentor feedback, submit well before the deadline. → Continue refining features, contributing to Octave community long-term. |

HISTORY OF MY CONTRIBUTIONS:

1. [\[statistics\] #173](#): Adds new functionalities to glmfit
2. [\[octave forge\] \(statistics\) #66388](#) anovan: suggested help text example not working
3. [#66595](#) The qz function documentation was not updated with generalized eigenvalues are no longer returned (**Commit Applied**)
4. [#66629](#) print pdf does not work
5. [#66650](#) barh with nonzero baseline...
6. [\[symbolic\] #1311](#): Modified Test Expectations
7. [\[symbolic\] #1316](#): Fix compatibility with BSD tar by supporting both tar and gtar

8. [\[symbolic\] #1319: Added the pol2cart functionality for Symbolic compatibility](#)
9. [\[symbolic\] #1317: warning\('on', 'all'\) leads to many Octave warnings](#)
10. [\[Discourse\] #6136: CIE1931 to spectrum plot](#)
11. [\[Discourse\] #6082: Need Matlab test of constant functions](#)

ACADEMIC EXPERIENCE AND OTHER ACTIVITIES:

I'm currently pursuing a **B.Tech in Computer Engineering** at **VJTI** (currently in Second Year), where I've developed **expertise** in **C++** and **Octave**, alongside a strong grasp of **clustering algorithms** and data structures. My academic projects have **sharpened** my skills in **Octave's classdef system**, which I've applied to design extensible statistical tools. As an **active GNU Octave contributor**, I've participated in Octave **Discourse discussions**, **assisted users** with troubleshooting, and **gained deep familiarity** with the statistics package codebase. **Solving algorithmic challenges** on [LeetCode](#) has further refined my **C++ problem-solving** abilities, ideal for performance-driven implementations. This technical proficiency, paired with **my passion** for open-source and **dedication to Octave**, **equips me to enhance** the statistics package **effectively through GSoC**. Hereby, I am also attaching my resume: [Swayam Shah](#).

